

#	Advanced Queries Overview	SQL concepts used
1	Hot-spot map: show the 10 parks / localities with the most distinct species and the 5 most distinct bird species	JOIN (Occurrence, Location), GROUP BY + aggregation (COUNT(DISTINCT ...))
2	Rare-bird alert feed: recent sightings whose species is statistically rare in that region	Correlated subquery that counts past-year occurrences, Outer query joins Taxon & Location
3	Top-100 birder leaderboard: rank users by the number of birds they have seen in descending order	3-way JOIN (User, Occurrence, Rarity), GROUP BY + COUNT
4	Species overlap explorer: birds seen in both Region A and Region B	Correlated subquery that finds the birds that are seen in 2 distinct regions

Database Connection

```

PS C:\Users\furka> mysql -h 34.27.66.196 -P 3306 -u root -p
Enter password: *****
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 332
Server version: 8.4.5-google (Google)

Copyright (c) 2000, 2025, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> SHOW DATABASES;
+-----+
| Database |
+-----+
| UrbanBird |
| information_schema |
| mysql |
| performance_schema |
| sys |
+-----+
5 rows in set (0.03 sec)

```

DDL Commands for Tables

RARITY:

```
CREATE TABLE RARITY (  
  rarity_id TINYINT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
  label VARCHAR(15) NOT NULL UNIQUE,  
  min_occ INT NOT NULL,  
  max_occ INT NOT NULL  
);
```

HABITAT:

```
CREATE TABLE HABITAT(  
  location_id VARCHAR(255),  
  species_id VARCHAR(255),  
  PRIMARY KEY(location_id,species_id),  
  FOREIGN KEY location_id REFERENCES LOCATION(location_id) ON DELETE CASCADE,  
  FOREIGN KEY species_id REFERENCES TAXON(species_id) ON DELETE CASCADE  
);
```

USER:

```
CREATE TABLE `UrbanBird`.`NewUser` (  
  user_id VARCHAR(200) PRIMARY KEY,  
  email VARCHAR(200),  
  password VARCHAR(200)  
);
```

OCCURENCE:

```
CREATE TABLE UrbanBird.OCCURRENCE (  
  occurrence_id BIGINT PRIMARY KEY,  
  species_id VARCHAR(200),  
  location_id VARCHAR(200),  
  user_id INT,  
  event_date DATE,  
  individual_count INT,  
  FOREIGN KEY (species_id) REFERENCES TAXON(species_id),  
  FOREIGN KEY (location_id) REFERENCES LOCATION(location_id),  
  FOREIGN KEY (user_id) REFERENCES USER(user_id)
```

```
);
```

LOCATION:

```
CREATE TABLE "UrbanBird"."LOCATION" (  
    location_id VARCHAR(200) PRIMARY KEY,  
    latitude DECIMAL(9,6),  
    longitude DECIMAL(9,6),  
    locality VARCHAR(200),  
    country_code CHAR(2)  
);
```

TAXON:

```
CREATE TABLE "UrbanBird"."TAXON" (  
    species_id INT PRIMARY KEY,  
    scientific_name VARCHAR(100),  
    common_name VARCHAR(100),  
    species_code CHAR(6),  
    category VARCHAR(20),  
    order_id INT,  
    rarity_id INT,  
    FOREIGN KEY (order_id) REFERENCES "ORDER"(order_id),  
    FOREIGN KEY (rarity_id) REFERENCES "RARITY"(rarity_id)  
);
```

ORDER:

```
CREATE TABLE "UrbanBird"."ORDER" (  
    order_id INT PRIMARY KEY,  
    order_name VARCHAR(40)  
);
```

3 Tables with at least 1000 records

```
1 SELECT COUNT(*) FROM LOCATION;
```

Results
COUNT(*)

```
1 SELECT COUNT(*) FROM OCCURRENCE;
```

Results
COUNT(*)

4892

```
1 SELECT COUNT(*) FROM TAXON;
```

Results

COUNT(*)

16753

Advanced Queries

1 Hot-spot map

Part 1:

```
SELECT
```

```
    occurrence_location_taxon.locality, occurrence_location_taxon.common_name,  
    occurrence_location_taxon.observation_count
```

```
FROM (
```

```
    SELECT
```

```
        loc.locality, tax.common_name, COUNT(DISTINCT occur.occurrence_id) AS
```

```
observation_count
```

```
    FROM UrbanBird.OCCURRENCE occur
```

```
    JOIN UrbanBird.LOCATION loc ON occur.location_id = loc.location_id
```

```
    JOIN UrbanBird.TAXON tax ON occur.species_id = tax.species_id
```

```
    GROUP BY loc.locality, tax.common_name
```

```
) AS occurrence_location_taxon
```

```
WHERE (
```

```
    SELECT COUNT(*)
```

```
    FROM (
```

```
        SELECT
```

```
            loc2.locality, tax2.common_name, COUNT(DISTINCT occur2.occurrence_id) AS
```

```
species_count
```

```
        FROM UrbanBird.OCCURRENCE occur2
```

```
        JOIN UrbanBird.LOCATION loc2 ON occur2.location_id = loc2.location_id
```

```
        JOIN UrbanBird.TAXON tax2 ON occur2.species_id = tax2.species_id
```

```

WHERE loc2.locality = occurance_location_taxon.locality
GROUP BY loc2.locality, tax2.common_name
HAVING COUNT(DISTINCT occur2.occurrence_id) >
occurance_location_taxon.observation_count
) AS occurance_location_taxon2
) < 5
ORDER BY occurance_location_taxon.locality, occurance_location_taxon.observation_count
DESC
LIMIT 50;

```

locality	common_name	observation_count
Anahuac National Wildlife Refuge (TX)	Fulvous Whistling-Duck	26
Anahuac National Wildlife Refuge (TX)	White-faced Ibis	26
Anahuac National Wildlife Refuge (TX)	Black-necked Stilt	25
Anahuac National Wildlife Refuge (TX)	Pied-billed Grebe	24
Anahuac National Wildlife Refuge (TX)	Eurasian Moorhen	23
Angelina National Forest (TX)	Northern Cardinal	1
Aransas (TX)	American White Pelican	11
Aransas (TX)	Whooping Crane	10
Aransas (TX)	Great Blue Heron	6
Aransas (TX)	Reddish Egret	6
Aransas (TX)	Caspian Tern	5
Aransas National Wildlife Refuge (TX)	Whooping Crane	20
Aransas National Wildlife Refuge (TX)	Great Blue Heron	15
Aransas National Wildlife Refuge (TX)	Pied-billed Grebe	15
Aransas National Wildlife Refuge (TX)	Turkey Vulture	10
Aransas National Wildlife Refuge (TX)	Belted Kingfisher	9
Aransas National Wildlife Refuge (TX)	Snowy Egret	9
Attwater Prairie Chicken National Wildlife Refuge (TX)	Savannah Sparrow	12
Attwater Prairie Chicken National Wildlife Refuge (TX)	Killdeer	8
Attwater Prairie Chicken National Wildlife Refuge (TX)	Crested Caracara	7

Rows per page: 20 ▾ 1 – 20 of 50 | < > >|

Part 2:

Baseline:

Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=7911..7911 rows=50 loops=1)

idx_A – locality:

Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5346..5346 rows=50 loops=1)

idx_B – common_name:

Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5123..5123 rows=50 loops=1)

idx_C – location_id, species_id:

Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5188..5188 rows=50 loops=1)

idx_A:

```
1 CREATE INDEX idx_A ON UrbanBird.LOCATION(locality);
2
3 EXPLAIN ANALYZE
4
5 SELECT
6     |   | occurrence_location_taxon.locality, occurrence_location_taxon.common_name,
7     |   | occurrence_location_taxon.observation_count
8 FROM (
9     |   | SELECT
10    |   | |   | loc.locality, tax.common_name, COUNT(DISTINCT occur.occurrence_id) AS observation_count
11    |   | |   | FROM UrbanBird.OCCURRENCE occur
12    |   | |   | JOIN UrbanBird.LOCATION loc ON occur.location_id = loc.location_id
13    |   | |   | JOIN UrbanBird.TAXON tax ON occur.species_id = tax.species_id
14    |   | |   | GROUP BY loc.locality, tax.common_name
```

Results

Export ▾

EXPLAIN

-> Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5346..5346 rows=50 loops=1) -> Sort: occurrence_location_...

idx_B:

```
1 CREATE INDEX idx_B ON UrbanBird.TAXON(common_name)
2
3 EXPLAIN ANALYZE
4
5 SELECT
6     |   | occurrence_location_taxon.locality, occurrence_location_taxon.common_name,
7     |   | occurrence_location_taxon.observation_count
8 FROM (
9     |   | SELECT
10    |   | |   | loc.locality, tax.common_name, COUNT(DISTINCT occur.occurrence_id) AS observation_count
11    |   | |   | FROM UrbanBird.OCCURRENCE occur
12    |   | |   | JOIN UrbanBird.LOCATION loc ON occur.location_id = loc.location_id
13    |   | |   | JOIN UrbanBird.TAXON tax ON occur.species_id = tax.species_id
14    |   | |   | GROUP BY loc.locality, tax.common_name
15 ) AS occurrence_location_taxon
16 WHERE (
```

Results

Export ▾

EXPLAIN

-> Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5123..5123 rows=50 loops=1) -> Sort: occurrence_location_...

idx_C:

```
1 CREATE INDEX idx_C ON UrbanBird.OCCURRENCE(location_id, species_id);
2
3 EXPLAIN ANALYZE
4
5 SELECT
6     occurrence_location_taxon.locality, occurrence_location_taxon.common_name,
7     occurrence_location_taxon.observation_count
8 FROM (
9     SELECT
10         loc.locality, tax.common_name, COUNT(DISTINCT occur.occurrence_id) AS observation_count
11     FROM UrbanBird.OCCURRENCE occur
12     JOIN UrbanBird.LOCATION loc ON occur.location_id = loc.location_id
13     JOIN UrbanBird.TAXON tax ON occur.species_id = tax.species_id
14     GROUP BY loc.locality, tax.common_name
```

Results Export

EXPLAIN

-> Limit: 50 row(s) (cost=2.6..2.6 rows=0) (actual time=5188..5188 rows=50 loops=1) -> Sort: occurrence_location_...

When running “EXPLAIN ANALYZE” on the SQL query - we achieved a baseline cost of approximately 2.6 with a baseline runtime of 7911 ms.

We tested 3 different indexing strategies, indexing on LOCATION(locality) - idx_A, indexing on TAXON(common_name) - idx_B, and indexing on OCCURRENCE(location_id, species_id) - idx_C. Unfortunately, the cost for all queries (baseline, idx_A, idx_B, idx_C) did not change - with a value of 2.6 for all. We think that the cost did not change since this query uses 2 derived tables (subqueries that join multiple tables together), which MySQL stores as temporary tables. We believe that the temporary tables are not taken into account when calculating cost - therefore, this would explain the non-changing cost value. However, the queries did decrease the runtime. This leads to the conclusion that indexing does not improve the cost performance of this query but does improve the runtime. Thus, we intend to use the indexing strategy described in idx_B, which has the best runtime performance.

2 Rare-bird alert

Part 1:

```
CREATE VIEW past_year_occ_stats AS
SELECT species_id, location_id, SUM(individual_count) AS past_year_cnt
FROM OCCURRENCE
WHERE event_date >= CURDATE() - INTERVAL 1 YEAR
GROUP BY species_id, location_id;
```

```

SELECT l.locality, pys.species_id, r.rarity_id
FROM past_year_occ_stats AS pys
JOIN TAXON AS t USING (species_id)
JOIN RARITY AS r USING (rarity_id)
JOIN LOCATION AS l USING (location_id)
WHERE r.label <> 'Common'
GROUP BY l.locality, pys.species_id, r.rarity_id
ORDER BY l.locality, pys.species_id
LIMIT 15;

```

Results Export

locality	species_id	rarity_id
Anahuac National Wildlife Refuge (TX)	Himantopus mexicanus	3
Anahuac National Wildlife Refuge (TX)	Plegadis chihi	3
Chambers (TX)	Himantopus mexicanus	3
Fort Bend (TX)	Nyctanassa violacea	2
Galveston (TX)	Cardinalis cardinalis	2
Galveston (TX)	Platalea ajaja	3
Harris (TX)	Anser cygnoides	1
Harris (TX)	Ardea alba	3
Harris (TX)	Ardea herodias	3
Harris (TX)	Buteo jamaicensis	1
Harris (TX)	Cairina moschata	1
Harris (TX)	Cardinalis cardinalis	2
Medina (TX)	Falco sparverius	2
Montgomery (TX)	Sitta pusilla	1
Nueces (TX)	Spatula discors	3

Rows per page: 20 1 - 15 of 15 |< < > >|

Part 2:

Baseline (no extra index):

Limit: 15 row(s) (cost=2491..2491 rows=15) (actual time=2.95..2.95 rows=15 loops=1)

idx_A - event_date:

Limit: 15 row(s) (cost=36.38 rows=5.25) (actual time=1.01..1.01 rows=15 loops=1)

idx_B - location_id, species_id:

Limit: 15 row(s) (cost=2491..2491 rows=15) (actual time=5.73..5.74 rows=15 loops=1)

idx_C - location_id, species_id, event_date:

Limit: 15 row(s) (cost=2491..2491 rows=15) (actual time=3.10..3.11 rows=15 loops=1)

Running EXPLAIN ANALYZE on the rare-bird query gave a baseline planner cost of = 2 491 (full scan of the temporary view).

Adding idx_A on event_date (CREATE INDEX idx_occ_event_date ON OCCURRENCE(event_date);) slashed the cost to 36.4, a > 98 % reduction, because the optimiser switches to a date-range scan that feeds far fewer rows into the grouping step. By contrast, idx_B on (location_id, species_id) and idx_C on (location_id, species_id, event_date) both left the cost unchanged at \approx 2 491: their leading columns are not selective, so MySQL ignored them, and in idx_C the date column's third position prevents a range scan.

We therefore keep idx_A only. Indexing event_date aligns with the guideline to target WHERE predicates; indexing location_id and species_id offered no benefit here even though they are valid join keys.

idx_A – event_date:

```
1 CREATE INDEX idx_occ_event_date
2   ON OCCURRENCE(event_date);
3
4 EXPLAIN ANALYZE
5 SELECT l.locality,
6        pys.species_id,
7        r.rarity_id
8 FROM   past_year_occ_stats AS pys
9 JOIN   TAXON      AS t USING (species_id)
10 JOIN  RARITY     AS r USING (rarity_id)
11 JOIN  LOCATION  AS l USING (location_id)
12 WHERE r.label <> 'Common'
13 GROUP BY l.locality, pys.species_id, r.rarity_id
14 LIMIT 15;
15
16 DROP INDEX idx_occ_event_date ON OCCURRENCE;
```

Results Export

EXPLAIN

-> Limit: 15 row(s) (cost=36..38 rows=5.25) (actual time=1..1.01 rows=15 loops=1) -> Table scan on <temporary> (cost=36..38 rows=5.25) (actual time=1..1 rows=15 loop...

idx_B – location_id, species_id:

```
1 CREATE INDEX idx_loc_spc_test
2   ON OCCURRENCE(location_id, species_id);
3
4 EXPLAIN ANALYZE
5 SELECT l.locality,
6        pys.species_id,
7        r.rarity_id
8 FROM   past_year_occ_stats AS pys
9 JOIN   TAXON      AS t USING (species_id)
10 JOIN  RARITY     AS r USING (rarity_id)
11 JOIN  LOCATION  AS l USING (location_id)
12 WHERE r.label <> 'Common'
13 GROUP BY l.locality, pys.species_id, r.rarity_id
14 LIMIT 15;
15
16 DROP INDEX idx_loc_spc_test ON OCCURRENCE;
```

Results Export

EXPLAIN

-> Limit: 15 row(s) (cost=2491..2491 rows=15) (actual time=5.73..5.74 rows=15 loops=1) -> Table scan on <temporary> (cost=2491..2498 rows=400) (actual time=5.73..5...

idx_C – location_id, species_id, event_date:

```
1 CREATE INDEX idx_loc_spc_date_test
2   ON OCCURRENCE (location_id, species_id, event_date);
3
4 EXPLAIN ANALYZE
5 SELECT l.locality,
6        pys.species_id,
7        r.rarity_id
8 FROM   past_year_occ_stats AS pys
9 JOIN   TAXON      AS t USING (species_id)
10 JOIN  RARITY     AS r USING (rarity_id)
11 JOIN  LOCATION  AS l USING (location_id)
12 WHERE r.label <> 'Common'
13 GROUP BY l.locality, pys.species_id, r.rarity_id
14 LIMIT 15;
15
16 DROP INDEX idx_loc_spc_date_test ON OCCURRENCE;
```

Results

Export

EXPLAIN

-> Limit: 15 row(s) (cost=2491..2491 rows=15) (actual time=3.1..3.11 rows=15 loops=1) -> Table scan on <temporary> (cost=2491..2498 rows=400) (actual time=3.1..3.1...

3 Top-100 birder leaderboard

Part 1

```
SELECT u.user_id, COUNT(*) AS total_sightings
FROM USER u
JOIN OCCURRENCE o USING (user_id)
JOIN TAXON t USING (species_id)
JOIN RARITY r USING (rarity_id)
GROUP BY u.user_id
ORDER BY total_sightings DESC
LIMIT 100;
```

Results

Export

user_id	total_sightings
User 120929	612
User 247179	434
User 1257	397
User 3712	254
User 40300	249
User 76090	246
User 86795	229
User 69501	166
User 201625	130
User 40150	124
User 41931	100
User 110790	100
User 45332	96
User 46215	95
User 10021	86
User 123214	69
User 4013	69
User 1587	67

Rows per page: 20 1 ~ 20 of 100 |< > |

PART 2

Cost without any indices = 34.4

```

1 CREATE INDEX idx_occurrence_user_id ON OCCURRENCE(user_id);
2
3 EXPLAIN ANALYZE
4 SELECT u.user_id, COUNT(*) AS total_sightings
5 FROM USER u
6 JOIN OCCURRENCE o USING (user_id)
7 JOIN TAXON t USING (species_id)
8 JOIN RARITY r USING (rarity_id)
9 GROUP BY u.user_id
10 ORDER BY total_sightings DESC;
11 |

```

Cost with index CREATE INDEX idx_occurrence_user_id ON OCCURRENCE(user_id) = 34.3

```

1 CREATE INDEX idx_occurrence_species_id ON OCCURRENCE(species_id);
2
3 EXPLAIN ANALYZE
4 SELECT u.user_id, COUNT(*) AS total_sightings
5 FROM USER u
6 JOIN OCCURRENCE o USING (user_id)
7 JOIN TAXON t USING (species_id)
8 JOIN RARITY r USING (rarity_id)
9 GROUP BY u.user_id
10 ORDER BY total_sightings DESC;
11

```

Cost with index CREATE INDEX idx_occurrence_species_id ON OCCURRENCE(species_id) = 34.4

```

1 CREATE INDEX idx_taxon_species_id ON TAXON(species_id);
2
3 EXPLAIN ANALYZE
4 SELECT u.user_id, COUNT(*) AS total_sightings
5 FROM USER u
6 JOIN OCCURRENCE o USING (user_id)
7 JOIN TAXON t USING (species_id)
8 JOIN RARITY r USING (rarity_id)
9 GROUP BY u.user_id
10 ORDER BY total_sightings DESC;
11

```

Cost with index CREATE INDEX idx_taxon_species_id ON TAXON(species_id) = 34.5

After trying out a few different indices, we came to the conclusion that indexing doesn't improve the cost performance of this query. We believe this is because when we create a primary or foreign key, mysql automatically create an index for that attribute; so, indexing something that has already been indexed doesn't cause major improvement in performance cost. In the context of this query, all the attributes in select, join, and groupby statements are either a part of primary

or foreign keys. Therefore, indexing does not help with improving the performance of this operation.

4 Species overlap explorer

PART 1

```
SELECT DISTINCT TAXON.common_name
FROM OCCURRENCE
JOIN LOCATION ON OCCURRENCE.location_id = LOCATION.location_id
JOIN TAXON ON OCCURRENCE.species_id = TAXON.species_id
WHERE LOCATION.locality = 'Harris (TX)'
AND OCCURRENCE.species_id IN (
  SELECT OCCURRENCE.species_id
  FROM OCCURRENCE
  JOIN LOCATION ON OCCURRENCE.location_id = LOCATION.location_id
  WHERE LOCATION.locality = 'Tarrant (TX)'
)
```

The screenshot shows the Google Cloud Platform console interface. At the top, there's a search bar and navigation links. Below, the 'All instances' section shows a database instance named 'mysql-team004-db'. The 'Editor 2' tab is active, displaying a SQL query. The query is a complex join between 'OCCURRENCE', 'LOCATION', and 'TAXON' tables, filtering for specific localities and species. The 'Run' button is highlighted. Below the query, the 'Results' section shows a table with one column, 'common_name', and 15 rows of bird species names. The table is scrollable. At the bottom right, there's a pagination control showing 'Rows per page: 20' and '1 - 15 of 15'.

common_name
Great-tailed Grackle
Red-shouldered Hawk
Eastern Phoebe
Grass Wren
Great Egret
Black Vulture
Cedar Waxwing
American Robin
Black-bellied Whistling-Duck
Wood Duck
Black-necked Stilt
Great Blue Heron
Little Blue Heron
Green Heron
Yellow-crowned Night-Heron

PART 2

No index: Cost = 90101

The screenshot shows a SQL IDE with two editor tabs. The active tab contains the following SQL code:

```
1 CREATE INDEX idx_species ON OCCURRENCE(location_id, species_id);
2
3 EXPLAIN ANALYZE
4
5 SELECT
6   DISTINCT TAXON.common_name
7 FROM
8   OCCURRENCE
9 JOIN
10  LOCATION
11 ON
12  OCCURRENCE.location_id = LOCATION.location_id
13 JOIN
14  TAXON
```

A syntax error message is displayed at the top right: "Syntax error at or near 'ANALYZE'". Below the code editor, the "Results" tab is active, showing the "EXPLAIN" output:

```
-> Table scan on <temporary> (cost=90101..93738 rows=290786) (actual time=89.9..89.9 rows=18 loops=1) -> Temporary table with deduplica...
```

The bottom of the results pane shows "Rows per page: 20" and "1 - 1 of 1".

C. Species taxon Index: Cost = 90082
Slight change.

The screenshot shows a SQL IDE with two editor tabs. The active tab contains the following SQL code:

```
1 CREATE INDEX idx_species_tax_id ON TAXON(species_id);
2
3 EXPLAIN ANALYZE
4
5 SELECT
6   DISTINCT TAXON.common_name
7 FROM
8   OCCURRENCE
9 JOIN
10  LOCATION
11 ON
12  OCCURRENCE.location_id = LOCATION.location_id
13 JOIN
14  TAXON
```

A syntax error message is displayed at the top right: "Syntax error at or near 'ANALYZE'". Below the code editor, the "Results" tab is active, showing the "EXPLAIN" output:

```
-> Table scan on <temporary> (cost=90082..93720 rows=290786) (actual time=106..106 rows=18 loops=1) -> Temporary table with deduplicati...
```

The bottom of the results pane shows "Rows per page: 20" and "1 - 1 of 1".

The final index design that we chose would be A, Locality Index. This index had a significant change in cost, going from 90101 to 9264. We believe that this is due to the reduction in search for the query. What we mean by that is without the index, the query has to search the whole table, which is over 5000 rows in order to find the correct localities. With this index, it reduces the search significantly

allowing for that huge decrease in cost. Regarding design B, Species Index, there was no change to the cost. We believe this is due to the query not requiring specific values for location_id and species_id, as the only time these are used are for the join resulting in this query to be equivalent to a full table search when joining. Lastly, design C, Species Index, the cost of this decreased slightly. This could have been due to slightly more efficient lookups to Taxon.species_id especially since taxon is a larger table and benefits more from having the index in place during the join.